

# CodeGen Agent – AgentCore POC Setup & Run Guide

## 1. Project Structure

```
Projects/
└── AI/
    ├── code_gen_agent.py
    ├── venv/          (Python env for development)
    └── vagent/        (Python env for AgentCore runtime)
```

## 2. Activate Environments

Activate development environment:

```
source venv/Scripts/activate
```

Activate AgentCore environment:

..

Here is the entire README as one single, clean, correctly formatted file.

This is **copy-paste ready** for GitHub — numbering works, spacing works, and formatting is polished.

# CodeGen Agent – AgentCore POC Setup & Run Guide

## 1. Project Structure

```
Projects/
└── AI/
    ├── code_gen_agent.py
    ├── venv/          (Python env for development)
    └── vagent/        (Python env for AgentCore runtime)
```

## 2. Activate Environments

Activate development environment:

```
source venv/Scripts/activate
```

Activate AgentCore environment:

```
source vagent/Scripts/activate
```

---

### 3. Configure & Launch the CodeGen Agent

Configure the agent:

```
agentcore configure -e code_gen_agent.py
```

Launch the runtime:

```
agentcore launch
```

---

### 4. Why We Switched from S3 Direct Code Deploy to Docker Container Deploy

Initially we attempted **Direct Code Deploy** (S3-based deployment).

However, this approach failed due to **package size limitations**.

#### Root Cause

The CodeGen Agent uses several heavy Python dependencies, including:

- boto3
- bedrock\_agentcore
- strands
- fastapi / pydantic
- supporting utilities

These made the deployment bundle exceed the size allowed for S3 direct deployment.

#### AgentCore Behavior

AgentCore automatically detected that the code bundle was too large and **fell back to container-based deployment**.

#### Why Docker Deployment Works

Docker containers overcome all the limitations because they:

- support large dependency sets
- can include native binaries
- can ship compiled wheels
- allow any project size
- provide a clean, reproducible environment

## Summary

S3 Direct Deploy fails due to size limits.

Docker deploy succeeds because it supports larger, dependency-heavy codebases.

---

## 5. Error Fixes During Setup

### Error 1 — Missing `uv`

**Warning:** Direct Code Deploy deployment unavailable (`uv` **not found**).

Fix:

```
pip install uv
```

---

### Error 2 — Missing `zip` Utility

**Warning:** Direct Code Deploy deployment unavailable (`zip` utility **not found**).

Install GnuWin32 Zip:

```
winget install GnuWin32.Zip
```

Add to PATH:

```
export PATH="$PATH:/c/PROGRA~2/GnuWin32/bin"
```

---

## 6. Test Payload

Use this to test the agent:

```
{  
  "requirement": "Create a simple FastAPI service with one GET /health endpoint."  
  "language": "python"  
}
```

## 7. Runtime Initialization Timeout Error

```
Unable to invoke endpoint successfully:  
Runtime initialization time exceeded.
```

This means the CodeGen Agent didn't complete startup within the 60-second initialization window.

### Common Causes

- Heavy imports at module load
- Large number of dependencies
- Slow initialization logic inside global scope
- Missing or misplaced entrypoint  
`(if __name__ == "__main__": app.run())`
- Unoptimized base image in the container

If you want, share your latest `code_gen_agent.py` and I'll help optimize startup time.